

# Raytracing


Adam Mally

CIS 561 Spring 2017

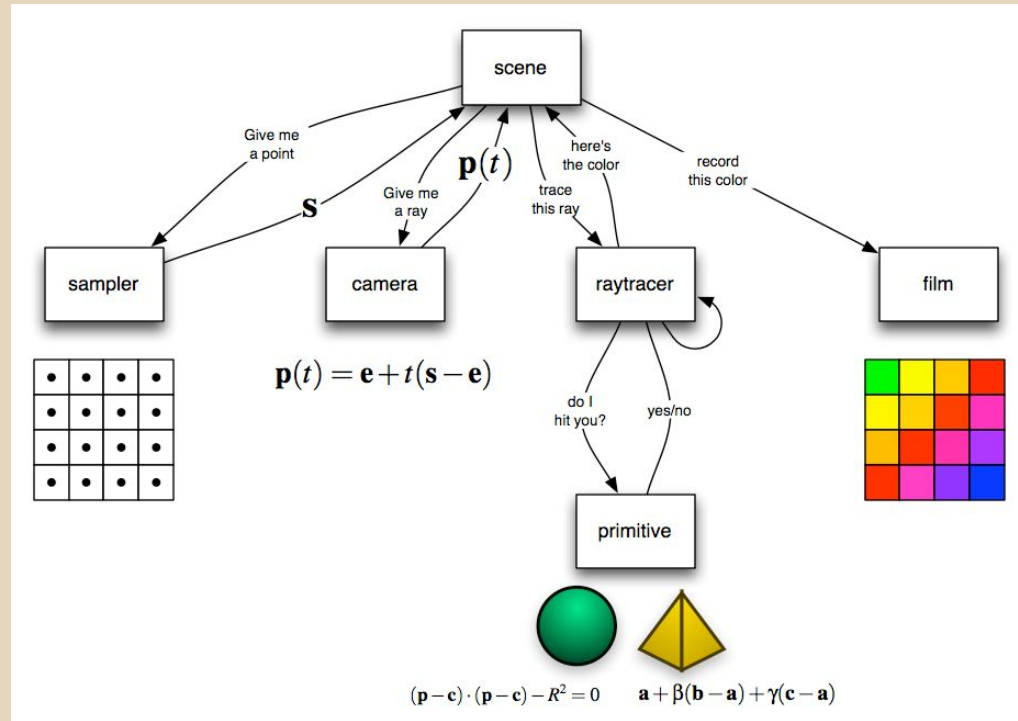
University of Pennsylvania

# What is raytracing?

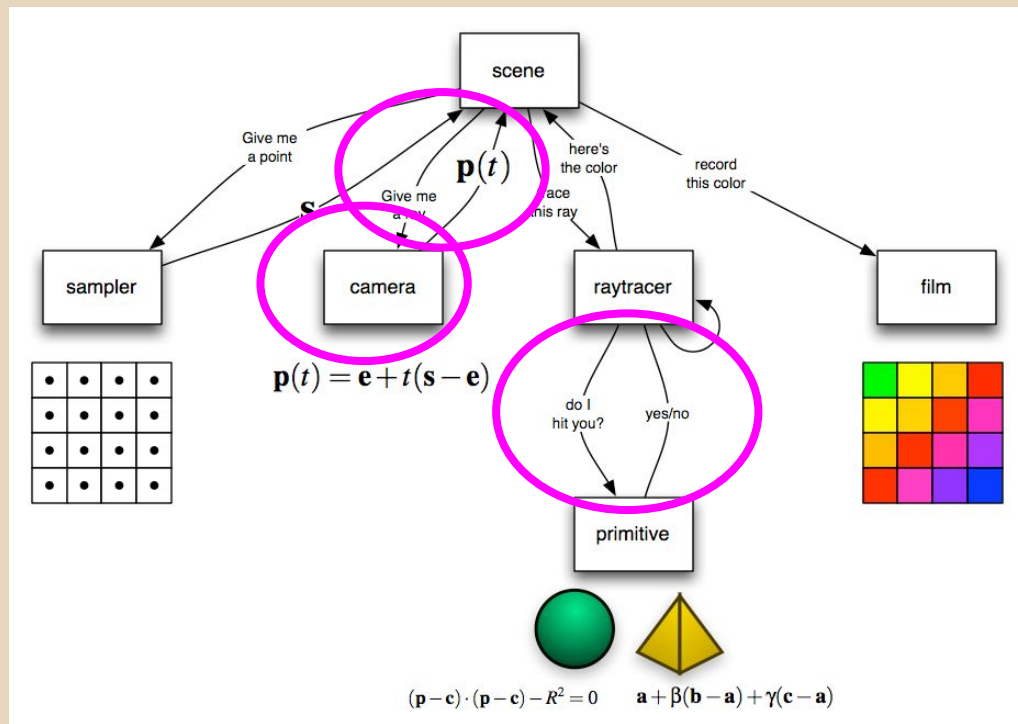
- Computer graphics technique used to simulate the way light bounces off surfaces
- Render fairly realistic-looking images without having to fully simulate the infinitely-many photons light sources give off
- Trace light paths backwards, from camera to light source(s)



# Basic Structure of a Raytracer




# Today's Topics




# Discussion: Camera Transformations

- What does it mean to transform a point from world space to camera space?
- What is the range of values for  $(S_x, S_y, S_z, S_w)$  and  $(P_x, P_y, P_z, P_w)$ ?
  - When  $C_z == \text{near\_clip}$ , what is  $P_z$ ?
  - When  $C_z == \text{far\_clip}$ , what is  $P_z$ ?
- How can we use this sequence of camera transformations to cast a ray that corresponds to a particular pixel in the screen?




# Review: View matrix parameters

- A space coordinate system is defined by two features: an origin and a set of axes
- **Camera origin:** simply the position the camera inhabits in world space
- **Forward direction (F):** Also known as the “look vector”. A direction in world coordinates that represents the direction in which the camera is looking. Represents the local Z-axis
- **Local right (R):** A direction in world coordinates that represents the direction that is “rightward” in the camera’s local coordinates. Represents local X-axis
  - Perpendicular to the **look direction**
- **Local up (U):** A direction in world coordinates that represents the direction that is “upward” in the camera’s local coordinates. Represents local Y-axis
  - Perpendicular to both **local right** and **look direction**




# Review: Perspective Frustum

- We can visualize the volume in which objects are visible to our camera
- Geometry outside the pyramid is not visible
- Shaped by several components:
  - Field of view
  - Aspect ratio: screen width / screen height
  - Near clip and far clip planes



# What is raycasting?


- Creating a line that passes through the viewing frustum and travels from the eye to some endpoint on a slice of the frustum (e.g. the far clip plane)
- The line's endpoint is determined by the pixel on our screen from which we want to raycast





# Normalized device coordinates


- Recall that your screen ranges from -1 to 1 on both the X and Y axes
- We can convert to NDC from any given pixel



- $sx = (2 * px / scr\_width) - 1$
- $sy = 1 - (2 * py / scr\_height)$
- px and py are the given pixel's x and y coordinates


# Screen point to world point

- $\text{ref} = \text{eye} + t * F$ 
  - $t = \text{any float} > 0$
- $\text{len} = |\text{ref} - \text{eye}|$
- $V = U * \text{len} * \tan(\alpha)$
- $H = R * \text{len} * \text{aspect} * \tan(\alpha)$
- $\alpha = \text{FOVY} / 2$
- $p = \text{ref} + s_x * H + s_y * V$ 
  - $s_x, s_y$  are in NDC



# Getting a ray from the world point

- $\text{ray\_origin} = \text{eye}$
- $\text{ray\_direction} = \text{normalize}(p - \text{eye})$
- Arbitrary point on ray =  $\text{eye} + t * \text{ray\_direction}$



# Faster method for computing a ray


- $P = \text{ViewMat}^{-1} * \text{ProjMat}^{-1} * ((px, py, 1, 1) * \text{farClip})$ 
  - $px, py$  are coordinates in NDC
  - $\text{ray\_origin} = \text{eye}$
  - $\text{ray\_direction} = \text{normalize}(P - \text{eye})$
- Pixel Z coord is technically arbitrary as long as it's in range (0, 1]
  - We use a value of 1 because we know where the far clip plane is, so un-homogenizing the vector is easy.
- What spatial transformation does this accomplish?
- Only need to compute the inverse of the view-projection matrix once per camera change

# What do we do with rays?

- Find their intersections with geometry in the scene
- Compute “fragment data” for these intersections
- Use this data to shade the pixels that correspond to each ray
- How do we find these intersections?


# Ray-polygon intersection

- Most common intersection test is ray-triangle
- We'll also cover ray-sphere and ray-cube
  - All three are commonly used in basic raytracer testing
- Generally, want to test for intersection against the **untransformed** geometry




# Frames of reference

- Before you try to test a ray against primitive geometry, you must first transform the ray so from its perspective, the geometry in question is primitive
- Simply transform the ray's direction and origin by the *inverse* of the geometry's model matrix



# Frames of reference


- Similarly, make sure you transform the results of your intersection test back into world space (e.g. the point of intersection, the surface normal at the intersection, etc.)





# Ray-sphere intersection

- Sphere defined as  $(x_s - x_c)^2 + (y_s - y_c)^2 + (z_s - z_c)^2 = r_s^2$ 
  - Sphere center =  $\langle x_c \ y_c \ z_c \rangle$
  - All points on the sphere surface =  $\langle x_s \ y_s \ z_s \rangle$
  - $r_s$  is the sphere's radius
- Ray defined as:  $R_0 + t * R_d$ 
  - $R_0 = \langle x_0 \ y_0 \ z_0 \rangle$
  - $R_d = \langle x_d \ y_d \ z_d \rangle$
  - $t$  is a parameterization of  $R_d$  (i.e. a float)



# Ray-sphere intersection

- Substitute  $\langle x_s y_s z_s \rangle$  for the ray equation:

$$(x_\theta + t*x_d - x_c)^2 + (y_\theta + t*y_d - y_c)^2 + (z_\theta + t*z_d - z_c)^2 = r_s^2$$

- Can also be written as:

- $At^2 + Bt + C = 0$

- $A = x_d^2 + y_d^2 + z_d^2$

- $B = 2(x_d(x_\theta - x_c) + y_d(y_\theta - y_c) + z_d(z_\theta - z_c))$

- $C = (x_\theta - x_c)^2 + (y_\theta - y_c)^2 + (z_\theta - z_c)^2 - r_s^2$

- Note that we now have a quadratic equation


- We can solve for  $t$  using the quadratic formula!

# Ray-sphere intersection

- $t_0, t_1 = (-B \pm \sqrt{B^2 - 4AC}) / (2A)$ 
  - $t_0$  is for the - case and  $t_1$  is for the + case
- Remember: if the discriminant is negative, then there is no real root and therefore no intersection
  - Discriminant =  $B^2 - 4AC$
- If  $t_0$  is positive, then we're done. If not, then compute  $t_1$ .


# Ray-sphere intersection

- Once we have  $t$ , we can plug it into our ray equation to find the closest point of intersection on our sphere.
  - If all we care about is whether or not we hit the sphere, we can just check:
    - $\text{near\_clip} < t < \text{far\_clip}$
- $P = R_0 + t * R_d$




# Ray-cube intersection

- Begin by storing  $t_{\text{near}} = -\text{infinity}$  and  $t_{\text{far}} = \text{infinity}$
  - For each pair of planes associated with the X, Y, and Z axes (the example uses the X “slab”):
    - If  $x_d$  is 0, then the ray is parallel to the X slab, so
      - If  $x_0 < x_{\min}$  or  $x_0 > x_{\max}$  then we miss
    - $t_0 = (x_{\min} - x_0) / x_d$
    - $t_1 = (x_{\max} - x_0) / x_d$
    - If  $t_0 > t_1$  then swap them
    - If  $t_0 > t_{\text{near}}$  then  $t_{\text{near}} = t_0$
    - If  $t_1 < t_{\text{far}}$  then  $t_{\text{far}} = t_1$
  - Repeat for Y and Z
  - If  $t_{\text{near}} > t_{\text{far}}$  then we miss the box
- 
- We can't see the  $z_{\text{slab}}$  because it's coming directly out of the screen
- $R_d = \langle x_d, y_d, z_d \rangle$
- $t_{\text{near}}$
- $t_{\text{far}}$
- $y_{\text{slab}}$
- $\langle x_{\max}, y_{\max} \rangle$




# Ray-plane intersection

- Plane defined as:  $\text{dot}(\mathbf{N}, (\mathbf{P} - \mathbf{S})) = 0$ 
  - $\mathbf{N}$  is the plane's normal
  - $\mathbf{S}$  is some point on the plane
  - $\mathbf{P}$  is the point of intersection
- Ray defined as:  $\mathbf{R}_0 + t * \mathbf{R}_d$
- Substitute P for ray:
  - $\text{dot}(\mathbf{N}, (\mathbf{R}_0 + t * \mathbf{R}_d - \mathbf{S})) = 0$
- Solve for t:
  - $t = \text{dot}(\mathbf{N}, (\mathbf{S} - \mathbf{R}_0)) / \text{dot}(\mathbf{N}, \mathbf{R}_d)$



# Point-in-triangle


- Use **barycentric coordinates** to test if  $P$  is within the bounds of a triangle
  - The **barycenter** of a triangle is its center of mass, often given unequal weighting to its vertices
- $S = \text{area}(P_1, P_2, P_3)$
- $S_1 = \text{area}(P, P_2, P_3)/S$
- $S_2 = \text{area}(P, P_3, P_1)/S$
- $S_3 = \text{area}(P, P_1, P_2)/S$
- Therefore,  $P = S_1P_1 + S_2P_2 + S_3P_3$
- So,  $P$  is within the triangle if:
  - $0 \leq S_1 \leq 1$
  - $0 \leq S_2 \leq 1$
  - $0 \leq S_3 \leq 1$
  - $S_1 + S_2 + S_3 = 1$



# Better triangle test

- Use only two weights to parameterize the triangle
- Point in triangle =  $(1 - u - v)P_1 + uP_2 + vP_3$ 
  - $u \geq 0, v \geq 0, (u + v) \leq 1$
- Point on ray =  $R_0 + t * R_d$
- Substitute ray equation in:
- $R_0 + t * R_d = (1 - u - v)P_1 + uP_2 + vP_3$
- Reformulate:

$$\circ \begin{bmatrix} -R_d, P_2 - P_1, P_3 - P_1 \end{bmatrix} \begin{vmatrix} t \\ u \\ v \end{vmatrix} = R_0 - P_1$$



“Fast, Minimum Storage Ray/Triangle Intersection” Moller and Trumbore

<https://www.cs.virginia.edu/~gfx/Courses/2003/ImageSynthesis/papers/Acceleration/Fast%20MinimumStorage%20RayTriangle%20Intersection.pdf>



# Better triangle test

- $$\begin{bmatrix} -R_d & P_2 - P_1 & P_3 - P_1 \end{bmatrix} \begin{vmatrix} t \\ u \\ v \end{vmatrix} = R_o - P_1$$
- Change to a format that can be solved with Cramer's Rule:
  - $P_2 - P_1 = E_1, P_3 - P_1 = E_2, R_o - P_1 = T$
  - $$\begin{vmatrix} t \\ u \\ v \end{vmatrix} = \frac{1}{\begin{vmatrix} -R_d & E_1 & E_2 \end{vmatrix}} \begin{vmatrix} \begin{vmatrix} T & E_1 & E_2 \end{vmatrix} \\ \begin{vmatrix} -D & T & E_2 \end{vmatrix} \\ \begin{vmatrix} -D & E_1 & T \end{vmatrix} \end{vmatrix}$$

# Better triangle test

$$\circ \quad \begin{vmatrix} t \\ u \\ v \end{vmatrix} = \frac{1}{\begin{vmatrix} -R_d & E_1 & E_2 \end{vmatrix}} \begin{vmatrix} \begin{vmatrix} T & E_1 & E_2 \end{vmatrix} \\ \begin{vmatrix} -D & T & E_2 \end{vmatrix} \\ \begin{vmatrix} -D & E_1 & T \end{vmatrix} \end{vmatrix}$$

- The length of a vector  $\langle A, B, C \rangle$  can also be expressed as  
 $-(A \times C) \cdot B = -(C \times B) \cdot A$

$$\begin{vmatrix} t \\ u \\ v \end{vmatrix} = \frac{1}{(R_d \times E_2) \cdot E_1} \begin{vmatrix} (T \times E_1) \cdot E_2 \\ (D \times E_2) \cdot T \\ (D \times E_1) \cdot D \end{vmatrix}$$

“Fast, Minimum Storage Ray/Triangle Intersection” Moller and Trumbore

<https://www.cs.virginia.edu/~gfx/Courses/2003/ImageSynthesis/papers/Acceleration/Fast%20MinimumStorage%20RayTriangle%20Intersection.pdf>

# Square intersection


- If we have a planar square in a scene, what methods could one use to find its intersection with a ray?

# Square intersection

- Find the point of intersection with the infinite plane the square lies in
- Test to see if this point lies within the bounds of the square
- Transform ray into square's local space and just test that  $-0.5 < P.x < 0.5$  and  $-0.5 < P.y < 0.5$ 
  - This assumes the unit square is aligned with the XY plane, is centered at the origin, and has a side length of 1.

# Square intersection

- If you know the vectors that make up the “width” and “height” of the square, you don’t even need to transform the ray into the square’s local space
  - If the following two statements are true, then the point-in-plane lies within the square
  - $\text{abs}(\text{dot}(\mathbf{P} - \text{square\_center}, \text{width\_vector})) < \text{length}(\text{width\_vector})^2$
  - $\text{abs}(\text{dot}(\mathbf{P} - \text{square\_center}, \text{height\_vector})) < \text{length}(\text{height\_vector})^2$
  - Why?



# Finding local normals

- When we've found an intersection in local object space, we usually want to find the local normal at that point
  - Normals are essential for most shading computations
- How can we find the object-space normal of a sphere?

# Finding local normals

- When we've found an intersection in local object space, we usually want to find the local normal at that point
  - Normals are essential for most shading computations
- How can we find the object-space normal of a sphere?
  - Normalize the point of intersection (assuming  $W$  is 0 beforehand)
- What about cubes?


# Finding local normals

- When we've found an intersection in local object space, we usually want to find the local normal at that point
  - Normals are essential for most shading computations
- How can we find the object-space normal of a sphere?
  - Normalize the point of intersection (assuming  $W$  is 0 beforehand)
- What about cubes?
  - Need to find which coordinate has the greatest magnitude
  - This determines which of the three major axes the normal lies along
  - Set the positivity of the normal by multiplying it with the **sign** of the largest-magnitude coordinate
  - Example: intersection is  $\langle -0.5, 0.45, -0.3 \rangle$ , so normal is  $\langle -1, 0, 0 \rangle$



# Surface normals

- Properly transforming surface normals from object space into world space is not as simple as multiplying them by the model matrix
  - Doing so skews them slightly, making them no longer normal (orthogonal) to the surface



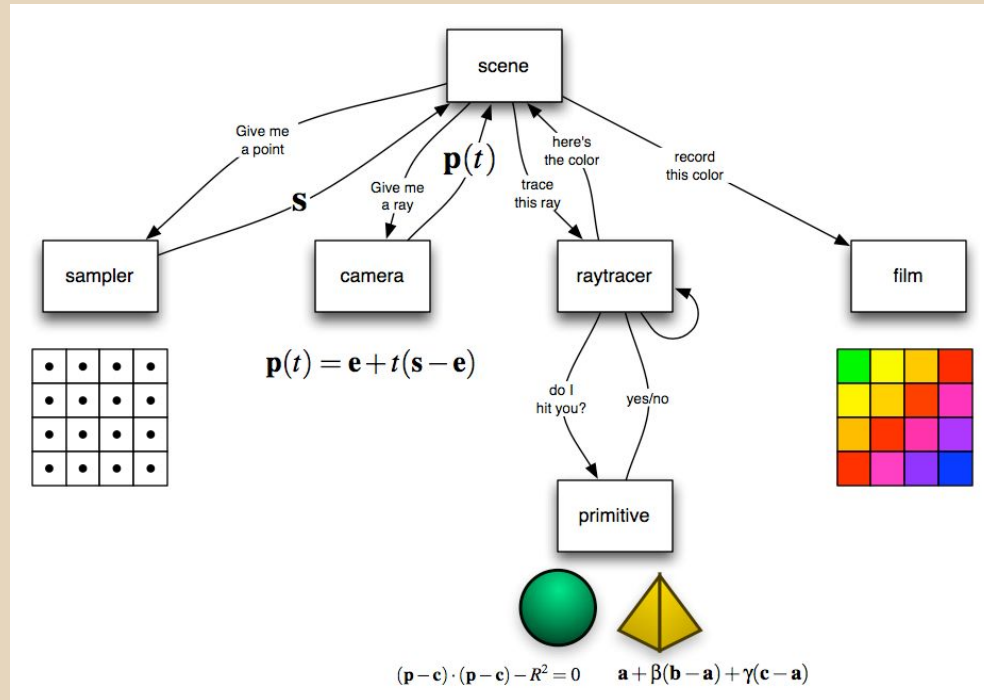
# Surface normals

- Of translate, rotate, and scale operations, only scale incorrectly transforms normals
- We want to **invert** the scale that is applied to the normal while keeping rotation the same
  - Translation has no effect on surface normal, so we ignore it entirely
- If we just invert the model matrix, both **rotation** and **scale** are inverted
- If we **invert and transpose** the model matrix, only **scale** is inverted
  - Inverting a rotation matrix  $R(\theta)$  is equivalent to making a matrix  $R(-\theta)$
  - Transposing a rotation matrix  $R(\theta)$  is also equivalent to making a matrix  $R(-\theta)$
  - Combining the inverse and transpose double-inverses the rotation, ultimately leaving it as just  $R(\theta)$

# Surface normals


- Given a point on a surface, there exists a surface normal  $n$  and some tangent vector  $t$ 
  - $\text{Dot}(n, t) = n^T t = 0$
- When the object is transformed by a model matrix  $M$ ,  $Mt = t'$ 
  - The transformed normal  $n'$  must remain orthogonal to  $t'$ , so we multiply  $n$  by some matrix  $S$  to get  $n'$
  - $0 = (n')^T t'$
  - $0 = (Sn)^T Mt$
  - $0 = (n)^T S^T Mt = n^T t$
  - The above identity implies that  $S^T M = I$ , which in turn implies that  $S^T = M^{-1}$ , therefore  $S = (M^{-1})^T$
- In summary, multiply the local-space surface normal by the *inverse transpose* of the model matrix to correctly bring it into world space
- Matrix **must not include** translation components before being inverted or it will transform the normals improperly
- This same proof can be found on page 94 of PBRT vol. 3

# Basic Structure of a Raytracer




# Backward Raytracing

- In the real world, photons are emitted by light sources and bounce off surfaces
- A small fraction of reflected photons reach our eyes, meaning the majority of emitted photons have no bearing on what we see
- To avoid computing the paths of all these extra photons, we trace the paths of photons in reverse
  - From each pixel of our camera screen into the scene, and back to the light source




# Basic Raytracing

- The most basic raytracing algorithm involves two main ray types:  
**camera rays** and **light-feeler rays**
- **Camera rays** are emitted through each pixel of the view screen and are tested against all geometry in the scene
- When a **camera ray** hits a scene object, a **light-feeler ray** is cast from the point of intersection to each light source in the scene
  - If a **light-feeler ray** reaches its light source then it contributes a portion of light to the overall color of the **camera ray** that created it
  - If a **light-feeler ray** is obstructed by an object in the scene, then it contributes pure black to the overall **camera ray** color





# Light-feeler issue: floating point error

- When computing the point of intersection with a surface, floating point error often results in the POI being slightly inside the object intersected
- When casting the light-feeler ray, it will intersect the object for which we are trying to compute the lighting
- This causes an effect commonly known as “shadow acne”



# Shadow acne: what to do?


- Ignore intersection with last object hit
  - Problem: prevents self-shadowing
- Use a minimum  $t$  value for the shadow test
  - Problem: A “good”  $t$  value depends on the scale of the scene and the distance the intersection is from the camera
- Use *doubles* instead of *floats* for extra precision
  - Problem: Increases memory needed to store the scene
- Offset the computed intersection by moving it along the surface normal
  - Should only move a \*very\* small amount (e.g. a factor above the smallest possible increment of a floating point,  $\sim 10^{-6}$ )
- PBRT chapter 3.9 discusses this at length






# Radiometry

- There are several important properties of light that make up the foundation of raytracing (excluding certain phenomena such as black holes)
  - A light ray travels in a perfectly straight line from one point to another
  - Light rays do not interfere with one another if they intersect
  - Given two points in space that can directly see one another, the amount of light emitted from point A towards point B is the same amount of light seen at point B (i.e. the light is invariant)




# Radiometry

- Within radiometry, there is a system of units and measures for illumination
- Take an optics (geometric) approach to light
- Again, treat light as if it travels in perfectly straight lines
  - Excellent approximation when the light's wavelength is much smaller than the objects with which the light interacts
  - Cannot model diffraction, interference, etc.



# Physics of lighting: Flux

- We want to simulate the gradual shading that objects exhibit in real life
- By using various shading models (Lambert, Phong, Gouraud, etc.) we can approximate the energy reflected at a surface point
- A light's intensity reduces as it travels away from its source
  - **Radiant flux** is the amount of energy passing through a region of space per unit of time, measured in Joules/sec (aka Watts) and commonly denoted as  $\Phi$
  - Both spheres surrounding the point light have the same amount of total **flux**, but any one local area of the larger sphere has less flux than a local area on the smaller sphere
    - Hence objects further from the light being more weakly illuminated




# Physics of lighting: Irradiance

- We can represent the amount of **flux** arriving at a surface as **irradiance** ( $E$ ), and the amount of **flux** leaving a surface as **radiant exitance** ( $M$ )
- Their unit of measurement is Watts/meter<sup>2</sup>
- The **irradiance equation** for our point light emission sphere is:

$$E = \Phi / (4\pi r^2)$$

- This is just flux/surface area
- The amount of energy (illumination) from this light falls off proportionally to the squared distance from the light since it becomes diffused over an increasingly larger area



# Physics of lighting: Lambert's law

- We can use the irradiance equation to help understand the origin of Lambert's law
  - Lambert's law: the amount of light arriving at a surface is proportional to the cosine of the angle between the light direction and surface normal
  - $E = \Phi \cos(\theta) / A_L$
- As  $\theta$  increases, the area of the lit surface increases, meaning the flux is distributed across a larger surface area, causing the irradiance to decrease for any one point on the lit surface
- Review:  $\text{dot}(A, B) = |A| |B| \cos(\theta)$ 
  - $\theta$  = angle between A and B
- Can re-write Lambert's equation as  $E = \Phi \text{dot}(N, L) / A_L$ 
  - Assuming N and L are normalized, of course
- The amount of light reaching a single point on a surface from a point light ( $A_L = 0$ ) can be represented simply as  $\text{clamp}(\text{dot}(N, L), \theta, 1)$

